

# E-commerce Platform System Design

## System Architecture

### Service Decomposition Strategy

I opted for a domain-driven microservices architecture over a monolith for the following reasons:

- **Scalability:** Independent scaling of services based on demand
- **Resilience:** Isolated failure domains prevent system-wide outages
- **Development velocity:** Smaller, focused teams can work independently
- **Technology flexibility:** Different services can use appropriate technologies

### Key Architecture Layers

#### 1. Client Layer

- Web Applications, Mobile Apps, Admin Portal, Third-party Integrations, API consumers
- Each client interacts through the API Gateway

#### 2. API Management Layer

- API Gateway: Entry point for all client requests, handling routing, composition, and protocol translation
- Rate Limiter: Prevents abuse and ensures fair resource allocation
- Auth Service: Centralizes authentication and authorization

#### 3. Core Services Layer

- User Service: Account management, preferences, permissions
- Product Service: Catalog management, pricing, availability
- Catalog Service: Categories, search indexes, product relationships
- Order Service: Order processing, status tracking
- Payment Service: Payment processing, refunds, financial records
- Cart Service: Shopping cart management
- Review Service: Product reviews and ratings
- Flash Sales Service: Limited-time promotion management
- Inventory Service: Stock tracking and availability
- Recommendation: Personalized product suggestions

#### 4. Infrastructure Services Layer

- Service Registry: Service discovery and registration
- Config Server: Centralized configuration management
- Logging Service: Aggregated logging and analysis
- Monitoring Service: Health checks and metrics collection
- Circuit Breaker: Fault tolerance and service isolation

#### 5. Asynchronous Processing Layer

- Message Queue: Reliable message delivery (Kafka/RabbitMQ)
- Event Bus: Real-time event distribution
- Task Scheduler: Delayed and scheduled tasks

#### 6. Data Storage Layer

- Relational Databases: Transactional data (Users, Orders, Payments)
- NoSQL Databases: Product catalog, reviews, flexible schema data
- Caching Layer: Multi-level caching strategy
- Data warehousing and analytics storage

## Database Design & Scaling

### Data Distribution Strategy

- **Relational data:** PostgreSQL for primary transactional data, with MySQL read replicas
- **Document data:** MongoDB for product catalog and reviews
- **Key-value data:** DynamoDB for session data and feature flags

### Scaling Strategies

- **Horizontal partitioning:** Sharding by user ID for user-specific data
- **Vertical partitioning:** Database-per-service pattern
- **Read replicas:** Scale read capacity independently
- **Connection pooling:** Efficient database connection management
- **CQRS pattern:** Separate read and write models for high-traffic services

### Data Consistency

- **Strong consistency:** For financial transactions and inventory updates
- **Eventual consistency:** For product catalog and user preferences
- **Transaction patterns:** Saga pattern for distributed transactions

## Caching Architecture

## Multi-layer Caching

- **Client-side cache:** Browser and mobile app caching
- **CDN cache:** Static assets, images, and product media
- **API Gateway cache:** API response caching
- **Distributed cache layer:** Redis clusters for session data, product catalog, user preferences

## Cache Policies

- **LRU eviction:** For general-purpose caches
- **TTL-based:** For time-sensitive data
- **Cache-aside pattern:** For database query results
- **Write-through:** For data that must be persisted
- **Cache invalidation:** Event-based cache clearing
- **Flash sales priority:** Special handling for high-demand events

## Asynchronous Processing

### Event-driven Architecture

- Message Queue for high-throughput events (Kafka)
- RabbitMQ for task management
- Consumer services process events asynchronously

### Use Cases

- Inventory updates
- Order processing
- Email notifications
- Analytics events
- User action tracking
- Real-time updates (flash sales)

## Fault Tolerance & High Availability

### Resilience Patterns

- Circuit breakers to prevent cascading failures
- Bulkhead pattern for resource isolation
- Retry with exponential backoff
- Fallback mechanisms for degraded functionality

### Service Health Management

- Health check APIs for all services
- Automated service restarts
- Self-healing deployment with Kubernetes

## **Multi-region Deployment**

### **Geographic Distribution**

- Primary and secondary regions with active-active configuration
- Cross-region replication for databases and caches
- Global load balancing and DNS services

### **Failover Strategy**

- Automated failover for critical services
- Data replication with minimal lag
- Global CDN with edge caching

## **API Rate Limiting**

### **Protection Mechanisms**

- Token bucket algorithm implementation
- User-based, IP-based, and service-based limits
- Burst handling for legitimate traffic spikes
- Distributed rate limit counters using Redis

## **Security Considerations**

### **Authentication & Authorization**

- OAuth 2.0 with JWT tokens
- Role-based access control
- Fine-grained permissions

### **Data Protection**

- Encryption in transit and at rest
- PCI DSS compliance for payment processing
- Data minimization and retention policies

### **API Security**

- Request validation and sanitization

- Protection against common attack vectors
- Monitoring for suspicious activities

## **Monitoring & Observability**

### **Observability Stack**

- Distributed tracing across services
- Centralized logging with correlation IDs
- Real-time metrics and dashboards
- Alerting based on business-critical thresholds

### **Performance Measurement**

- Key performance indicators (KPIs) tracking
- User experience metrics
- Business impact metrics (conversion rate, cart abandonment)